



Cosmos DB Performance Tuning

Mike Benkovich
Microsoft Azure MVP
mike@benko.com | www.benkoTIPS.com | @mbenko





Mike Benkovich

- Developer
- Cloud **Architect & Consultant**
- Live in **Minneapolis**
- Founder of **Imagine Technologies, Inc.**
- Developing Courses for **LinkedIn Learning**
- Blog www.benkoTIPS.com
- Follow **@mbenko** on **Twitter**
- Send me **Feedback!** mike@benkotips.com



Mike Benkovich

Enterprise Cloud Architect,
Consultant, Developer Tools Ev...



My Sessions ... this week!

- ❖ Permit to Cloud – Land with Confidence in Azure
Tuesday 6/7 11:00 am – Discovery D
- ❖ Performance Tuning Strategies for Cosmos DB
Tuesday 6/7 4:00 pm – Imagination B
- ❖ Infrastructure as Code Bake-off ARM vs Bicep vs TF
Wednesday 6/8 4:00 pm – Discovery B

**What we'll cover
today**

Why Cosmos
Provisioning Cosmos
Working with Data
Data Modeling Strategies
Performance and tuning

Cosmos DB Overview

A Simple **Global** Assigned NoSQL Database that can be replicated across regions with a click of a button

High **Scalability** and **Availability**

Guaranteed **low latency** and **provisioned** performance

Multi-**model** Multi-**API**

Configurable **Consistency** Options

Containers and Documents

- Database as collection of Documents in Containers
- Provisioned performance – RU's
- Stored as JSON Document
 - Document Explorer (portal)
 - Storage Explorer
- Supports multiple Record Types in Containers
- No referential rules or constraints
- No enforcement of data types

```
{  
  "id": "a40d0078",  
  "Title": "Add stuff to the Queue",  
  "Description": null,  
  "Type": "Chore",  
  "Owner": null,  
  "IsDone": false,  
  "createdAt": "2019-09-05T11:40:12.4466765Z",  
  "completeAt": null,  
  "_rid": "vr0UAPQ9W4sDAAAAAAA==",  
  "_self": "dbs/vr0UAA==/colls/...",  
  "_etag": "\\"38001cae-0000-0300-0000-5d70...\"",  
  "_attachments": "attachments/",  
  "_ts": 1567683617  
}
```

Some Terminology

RELATIONAL

- Database
- Table(s), View(s)
- Row
- Column
- Foreign Key
- Join
- Sharding Key

COSMOS DB (SQL API)

- ❖ Account/Database
- ❖ Container/Collection
- ❖ Item/Document (JSON)
- ❖ Property
- ❖ Reference
- ❖ Embedded Document
- ❖ Partition Key

What we'll cover
today

Why Cosmos
Provisioning Cosmos
Code patterns & tools
Data Modeling Strategies
Performance and tuning

Provisioning Cosmos

Selecting the API : SQL, Mongo, Cassandra, Gremlin, Table

Capacity model : Provisioned vs Serverless

Geo-Redundancy and Multi-region Writes

Networking

Backup

Encryption

Get Started with Cosmos

Microsoft Azure

Home > Azure Cosmos DB > bnk-demos-cosmos-eus

bnk-demos-cosmos-eus | Data Explorer

Overview Activity log Access control (IAM) Tags Diagnose and solve problems Quick start Notifications Data Explorer Settings Features Replicate data globally Default consistency Backup & Restore Firewall and virtual networks Private Endpoint Connections CORS Keys Add Azure Cognitive Search Add Azure Function Advanced security (preview) Locks Containers Browse

SQL API

DATA

CosmosBlog-v1, CosmosBlog-v2, CosmosBlog-v3, CovidData, msdn, blogposts, events, usergroups, NOTEBOOKS, Gallery, My Notebooks

Items

Items

id /BlogID

1 efeac264-e00f... 13c224af-3097-413...
2 890f2738-15d... 27604f05-86ad-47...
3 f764a791-4afe... 13c224af-3097-413...
4 acb1c288-65c6... 13c224af-3097-413...
5 e7fa0ec7-2b4e... 13c224af-3097-413...
6 862224e3-3d7... 13c224af-3097-413...
7 d6ddf4dc-312... 27604f05-86ad-47...
8 b607d734-4ba... 27604f05-86ad-47...
9 dc286420-272... 27604f05-86ad-47...
10 17643544-25c... 27604f05-86ad-47...
11 db4e54f0-2bfa... 27604f05-86ad-47...
12 acecf47b-3d0e... 27604f05-86ad-47...
13 a362202a-a7b... 27604f05-86ad-47...
14 4c75f8a6-4ec7... 27604f05-86ad-47...
15 93cfc7de-313d... 27604f05-86ad-47...
16 768eb870-622... 27604f05-86ad-47...
17 83a50a8c-a63... 27604f05-86ad-47...
18 9e2893c3-667... 27604f05-86ad-47...
19 1b90a404-1bf... 27604f05-86ad-47...
20 Load more

1 "PostRowID": 45,
2 "BlogID": "13c224af-3097-4132-915c-4ac0f01a3dbe",
3 "id": "e7fa0ec7-2b4e-4af8-1ae-43c2669f6c7b",
4 "Title": "TechMasters Captured the Elusive Gilded Ga",
5 "Description": "<p>This morning, club President Vin",
6 "PostContent": "<p>this morning, club President Vin",
7 "DateCreated": "2016-02-06T16:55:06.0000000",
8 "DateModified": "2016-02-18T14:36:38.9030000",
9 "Author": "Vince Bullinger",
10 "IsPublished": false,
11 "IsCommentEnabled": true,
12 "Raters": 0,
13 "Rating": 0,
14 "Slug": "TechMasters-Captured-the-Elusive-Gilded-Ga",
15 "IsDeleted": false,
16 "_rid": "isoEAJLExv8FAAAAAAAA==",
17 "_self": "dbs/isoEAA=/colls/isoEAJLExv8=/docs/isoE",
18 "_etag": "\"0b000d09-0000-0100-0000-60031ff70000\"",
19 "_attachments": "attachments/",
20 "_ts": 1610817527

What we'll cover
today

Why Cosmos
Provisioning Cosmos
Working with Data
Data Modeling Strategies
Performance and tuning

Tools to work with Cosmos Data

Azure Portal

Data Explorer

Jupyter Notebooks

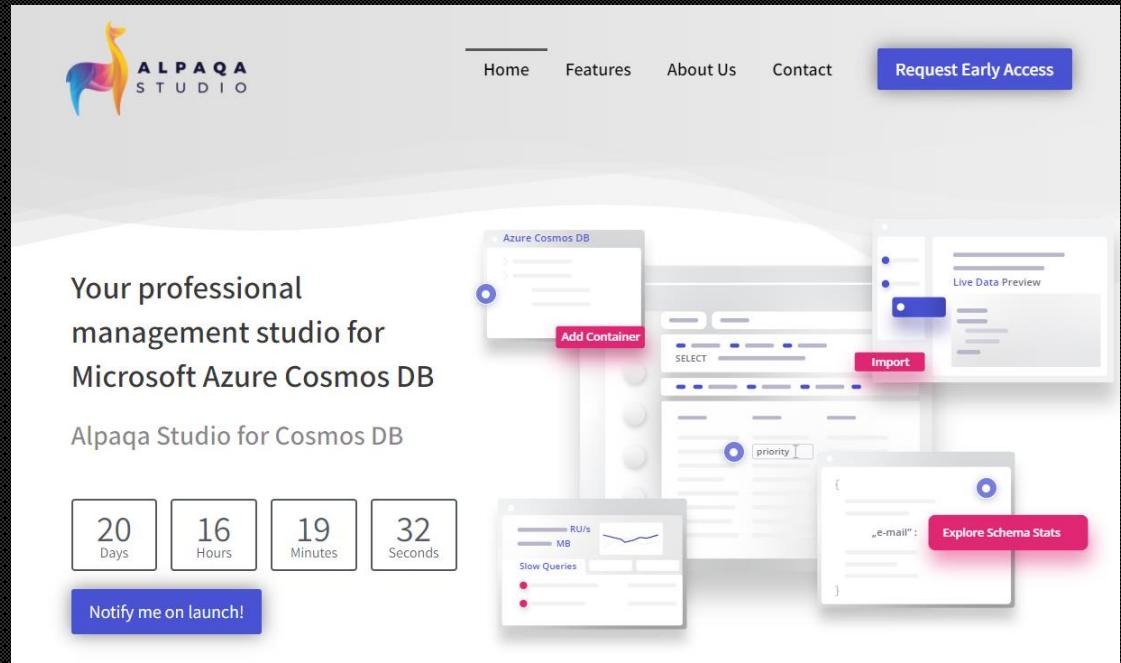
Alpaqa Studio

Editing Data (in place and bulk)

Explore Schema

Background task engine

SDK – Code



The screenshot shows the Alpaqa Studio website homepage. At the top right, there are navigation links for Home, Features, About Us, Contact, and a blue 'Request Early Access' button. The main content area features the Alpaqa Studio logo (a stylized orange and yellow flame-like icon) and the text: 'Your professional management studio for Microsoft Azure Cosmos DB'. Below this, there's a section for 'Alpaqa Studio for Cosmos DB' with four time-related boxes: '20 Days', '16 Hours', '19 Minutes', and '32 Seconds'. A large button at the bottom says 'Notify me on launch!'. To the right, there's a collage of interface snippets showing various tools: 'Azure Cosmos DB' (with 'Add Container' and 'Import' buttons), 'SELECT' queries, 'Live Data Preview', 'Slow Queries' monitoring, and 'Explore Schema Stats'.

<https://alpaqastudio.com>

[Alpaqa Studio - The IDE for Cosmos DB \(preview\)](#)

What we'll cover
today

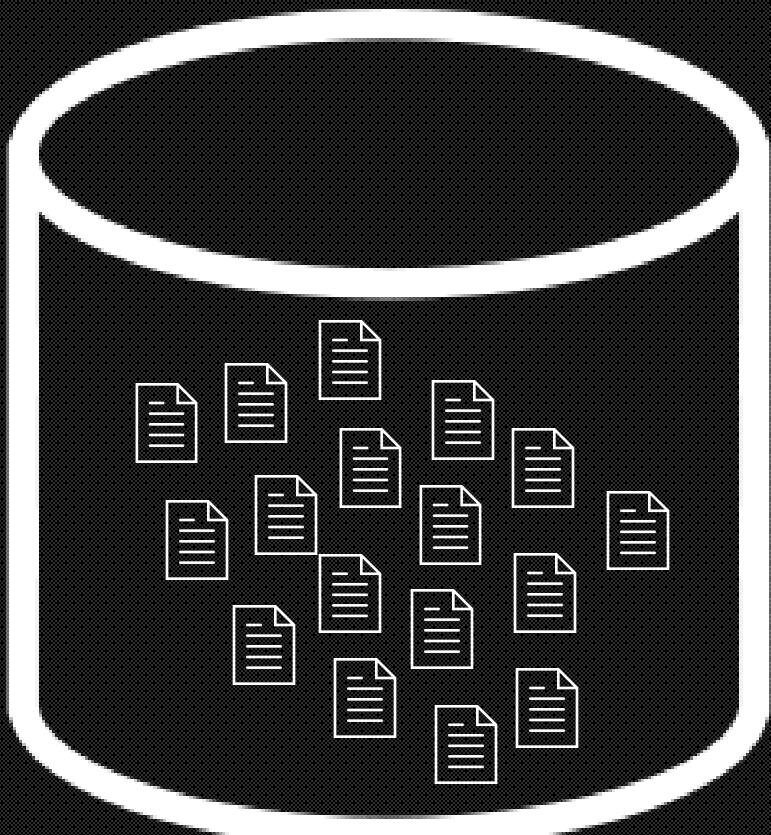
Why Cosmos
Provisioning Cosmos
Working with Data
Data Modeling Strategies
Performance and tuning

Provisioned Performance

- Expected performance has to be **provisioned**
- Expressed in **Request Units per second** (RU/s)
 - Represents the "cost" of a request in terms of CPU, memory and I/O
- Performance can be provisioned:
 - at the database-level
 - at the collection-level
- Provisioned performance **can be changed programmatically** with API calls

The rules of the game: partitioning

- Cosmos DB achieves horizontal scalability by **partitioning** your data



The rules of the game: partitioning

- Cosmos DB achieves horizontal scalability by **partitioning** your data

'userId': 'abc'



'userId': 'def'



'userId': 'ghi'



logical partitions

- Logical partitions group your data according to the collection's **partition key**



The rules of the game: partitioning

- A good partition key ensures **well-balanced partitions**
 - Both in terms of storage and in terms of throughput
- Ideally, read queries should get all their results from one partition

'userId': 'abc'



'userId': 'def'

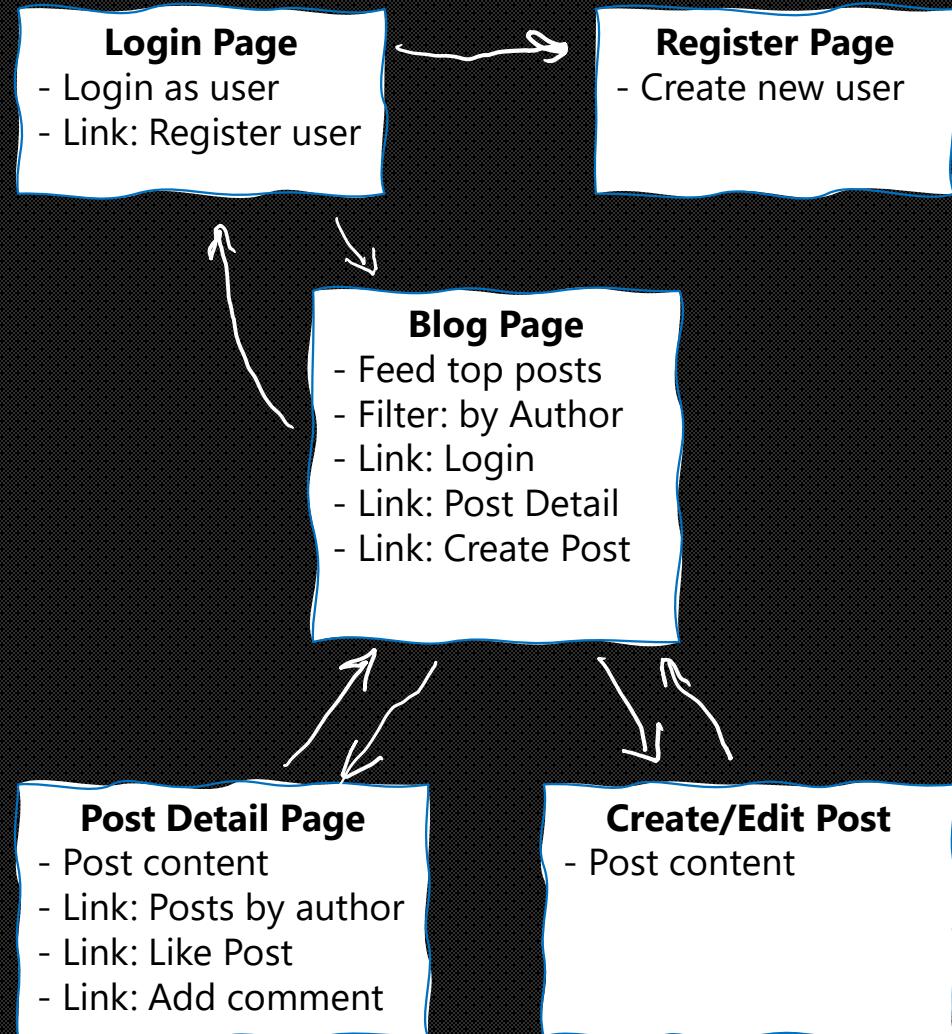


'userId': 'ghi'



```
SELECT * FROM u WHERE u.userId = 'def'
```

Scenario: Rehost Blog in Cosmos



Users can create **posts**. Users can also **like** and add **comments** to posts.

A front page displays a feed of recently created posts. It is possible to list all posts for a particular user. It is also possible to list all comments of a post and all users who have liked a post.

A post is displayed with its author's username and a count of comments and likes. Comments and likes are also displayed with the username of their authors.

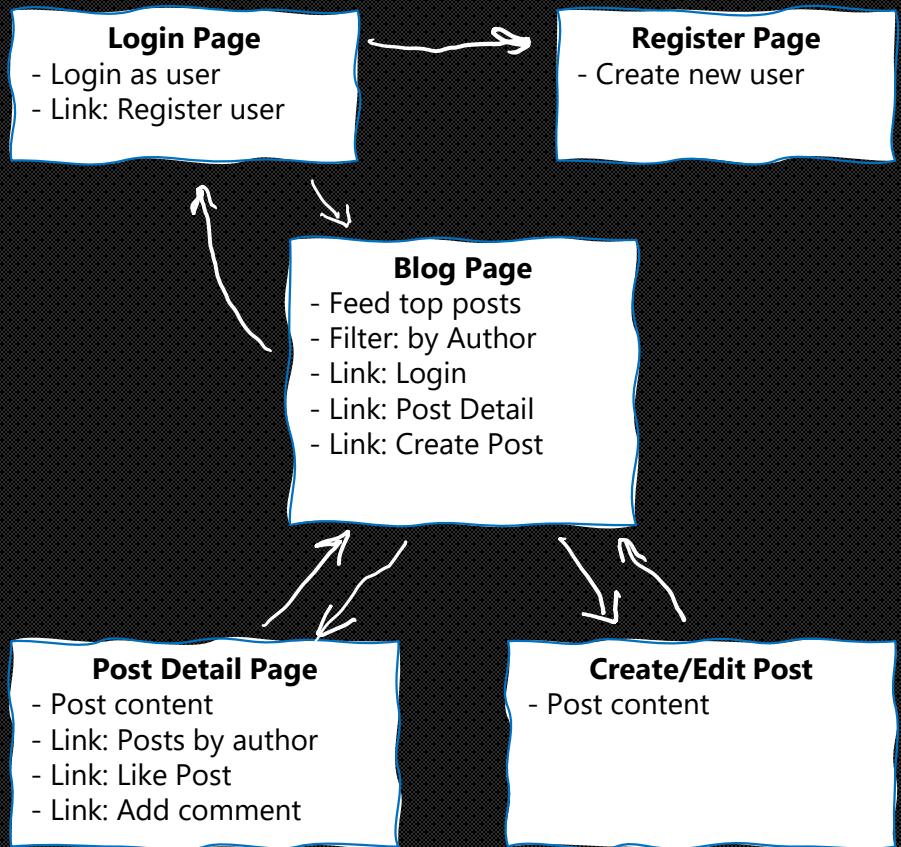
When displayed as lists, posts only present a truncated summary of their content.

Start by identifying the access patterns

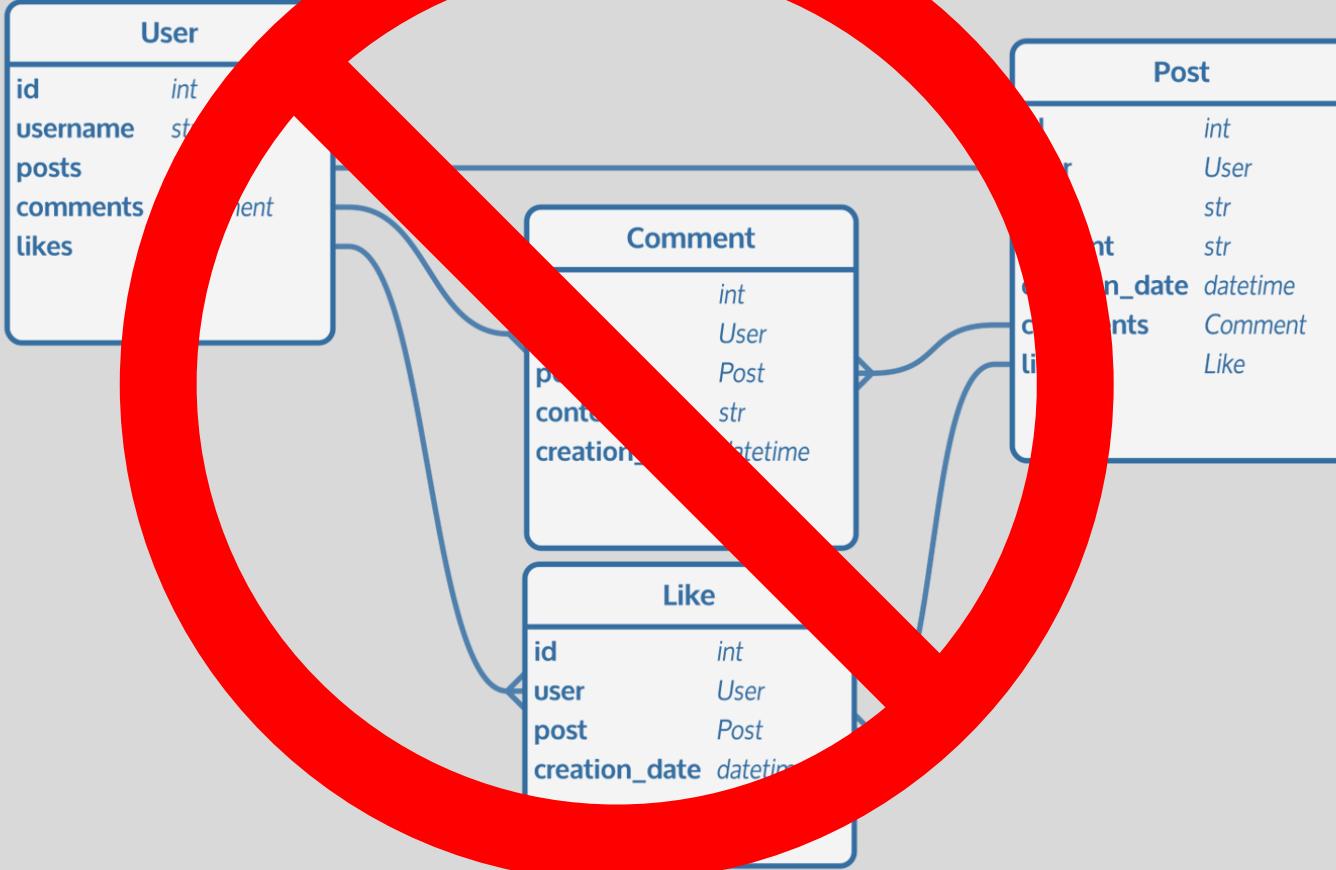
- What are the requests that our design will have to serve?
- To make it easier to follow, categorize them as:
 - **Command** (write request)
 - **Query** (read-only request)
- Think in terms of data navigation

Start by identifying the access patterns

- [C1] Create/edit a user's profile
- [Q1] Retrieve a user's profile
- [C2] Create/edit a post
- [Q2] Retrieve a post
- [Q3] List a user's posts in short form
- [C3] Create a comment
- [Q4] List a post's comments
- [C4] Like a post
- [Q5] List a post's likes
- [Q6] List the x most recent posts in short form (feed)



Maybe you instantly think about...



The rules of the game: data model

- In a Cosmos DB database,
 - like rows?



The rules of the game: data model

- In a Cosmos DB database, we store **documents**  in **collections** 



user



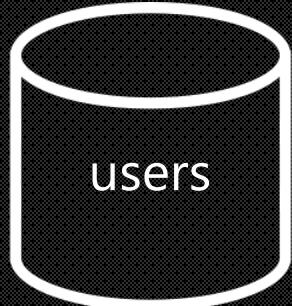
post



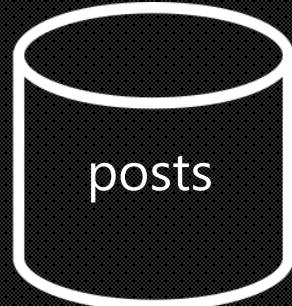
comment



like



users



posts

V1: a first shot

```
{  
  "id": "<user-id>",  
  "username": "<user-username>"  
}
```

```
{  
  "id": "<post-id>",  
  "userId": "<post-author-id>",  
  "title": "<post-title>",  
  "content": "<post-content>",  
  "creationDate": "<post-creation-date>"  
}
```

```
{  
  "id": "<comment-id>",  
  "postId": "<post-id>",  
  "userId": "<comment-author-id>",  
  "content": "<comment-content>",  
  "creationDate": "<comment-creation-date>"  
}
```

```
{  
  "id": "<like-id>",  
  "postId": "<post-id>",  
  "userId": "<liker-id>",  
  "creationDate": "<like-creation-date>"  
}
```

To embed or to reference?

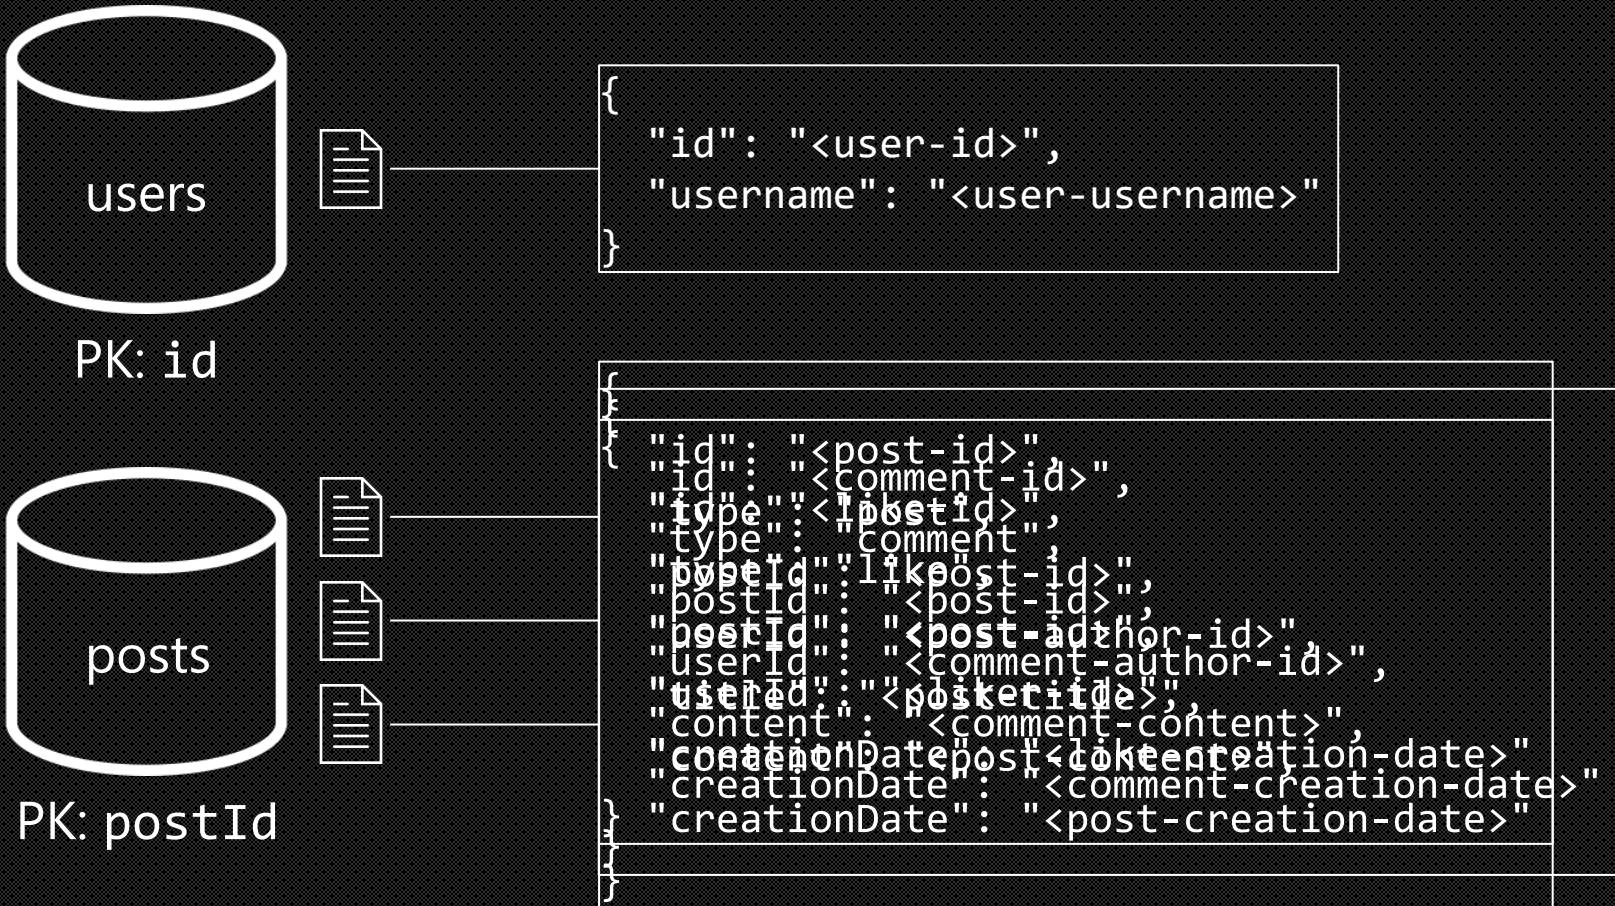
Embed when:

- 1:1
- 1:few
- Items queried together
- Items updated together

Reference when:

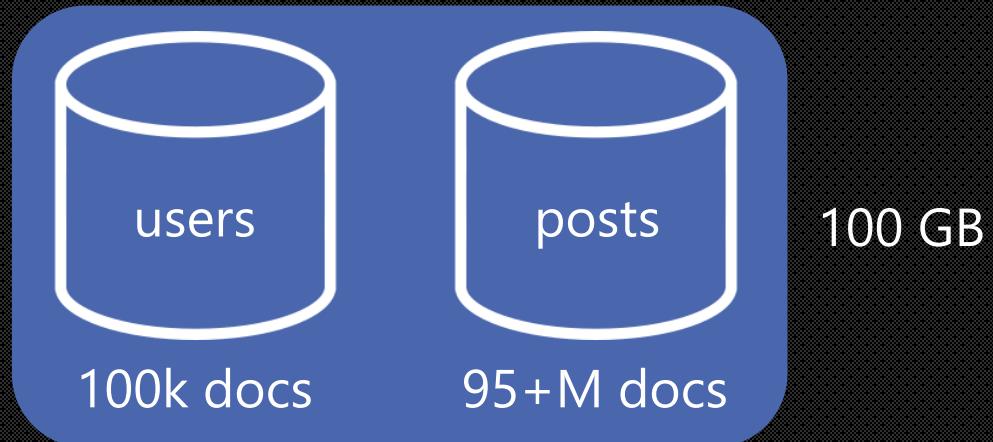
- 1:many (unbounded)
- Many:many
- Items updated independently

V1: a first shot



How well does our model perform?

- Benchmarked against a dummy dataset
 - 100,000 users
 - 5 – 50 posts / user
 - 0 – 25 comments / post
 - 0 – 100 likes / post



Create/edit a user's profile

C1 ✓

Q1

C2

Q2

Q3

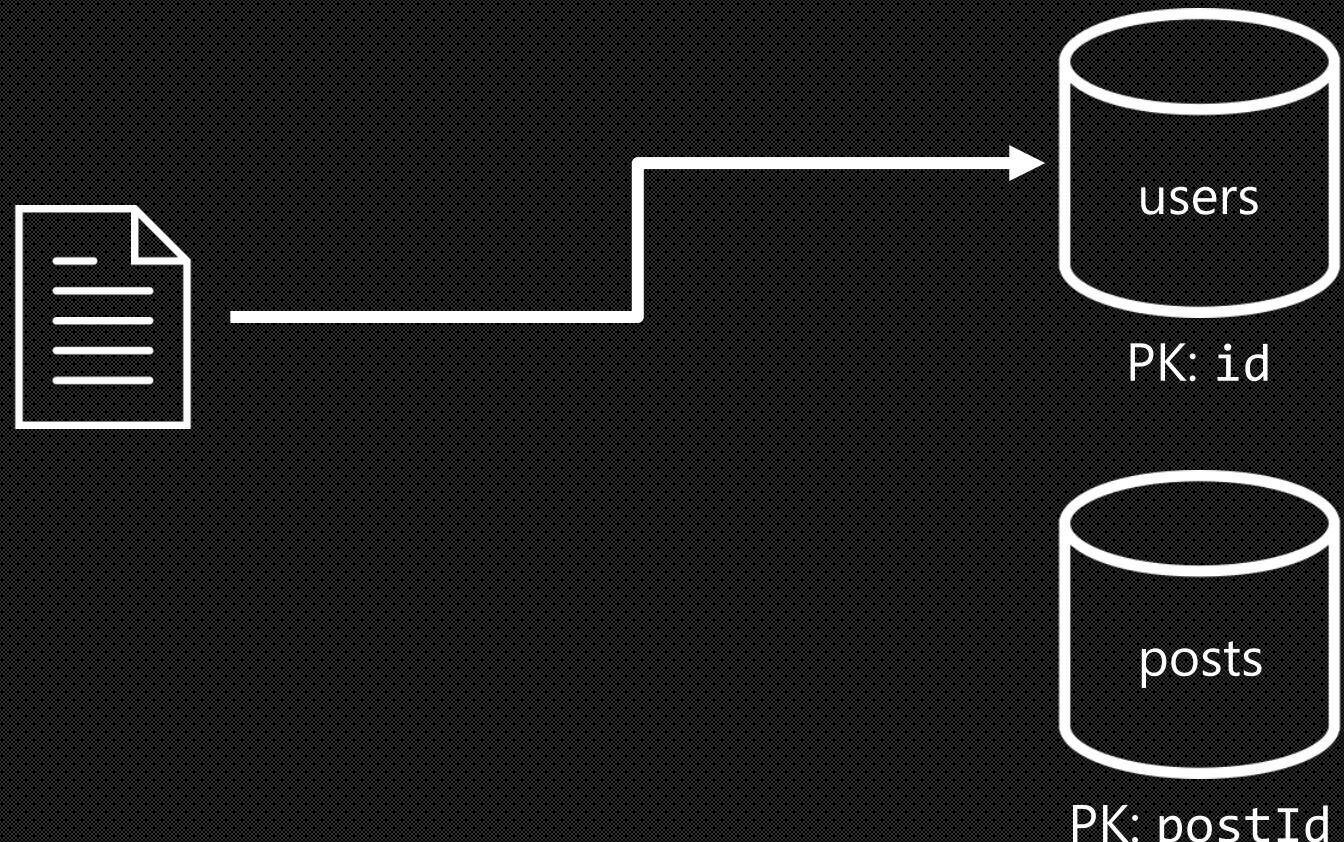
C3

Q4

C4

Q5

Q6



7 ms / 5.71 RU

Retrieve a user's profile

C1



Q1



C2

Q2

Q3

C3

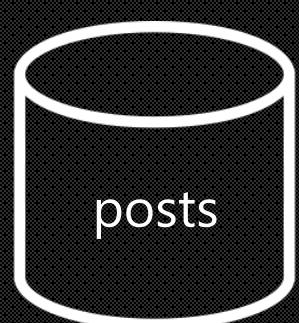
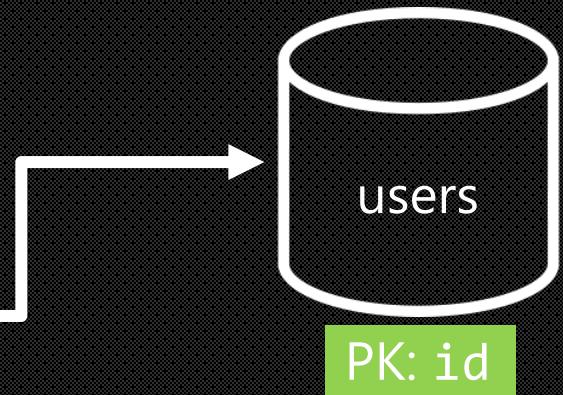
Q4

C4

Q5

Q6

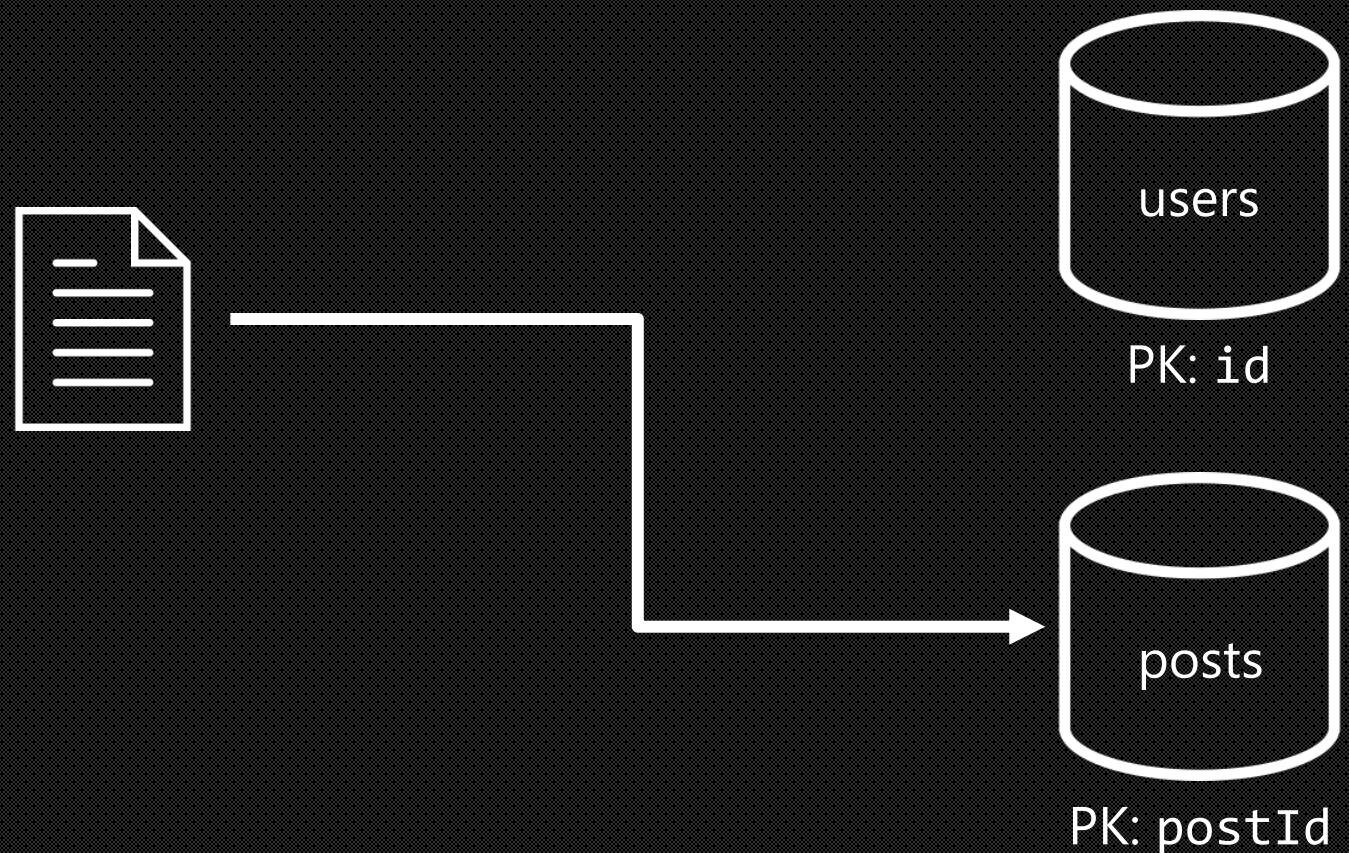
```
SELECT * FROM u WHERE u.id = '<user-id>'
```



3 ms / 2.90 RU

Create/edit a post

C1	✓
Q1	✓
C2	✓
Q2	
Q3	
C3	
Q4	
C4	
Q5	
Q6	



9 ms / 8.76 RU

Retrieve a post

C1 ✓

```
SELECT u.username FROM u  
WHERE u.id = '<post-author-id>'
```

Q1 ✓

```
SELECT * FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'post'
```

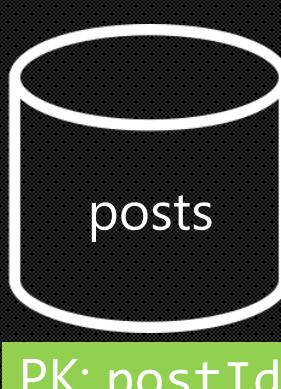
C2 ✓

```
SELECT COUNT(1) FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'comment'
```

Q2 !

```
SELECT COUNT(1) FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'like'
```

Q3



9 ms / 19.54 RU

List a user's posts in short form

C1



```
SELECT u.username FROM u  
WHERE u.id = '<post-author-id>'
```

Q1



```
SELECT * FROM p  
WHERE p.userId = '<author-id>'  
AND p.type = 'post'
```

C2



Q2



```
SELECT COUNT(1) FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'comment'
```

Q3



```
SELECT COUNT(1) FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'like'
```

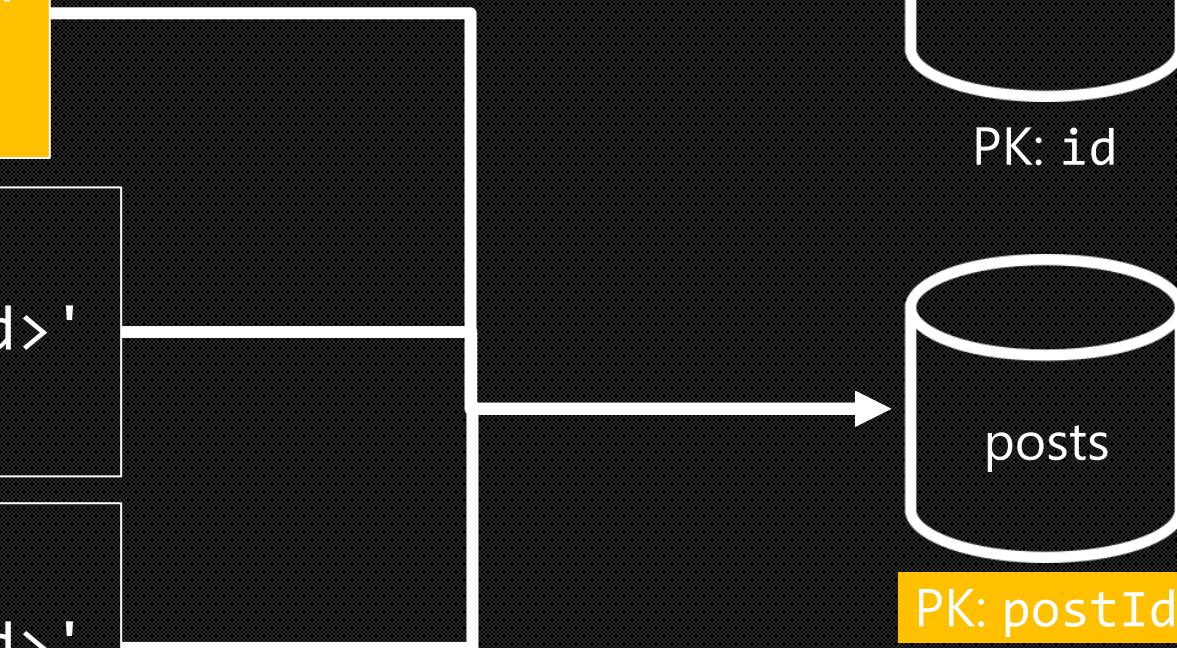
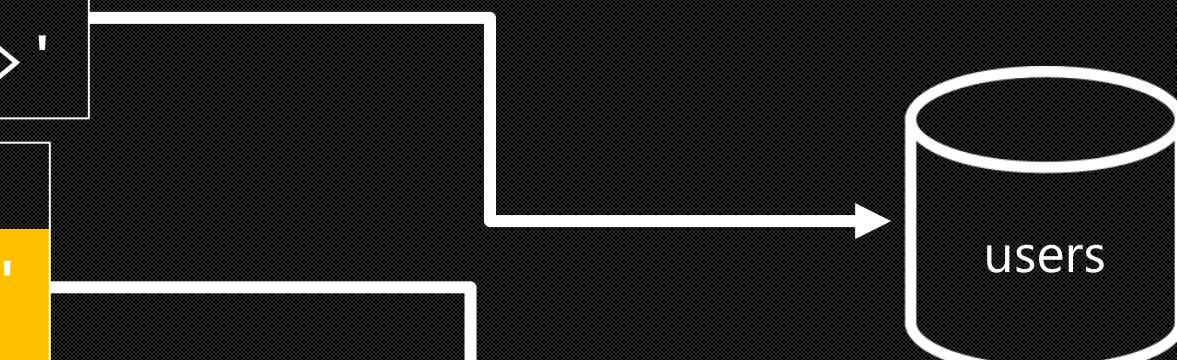
C3

Q4

C4

Q5

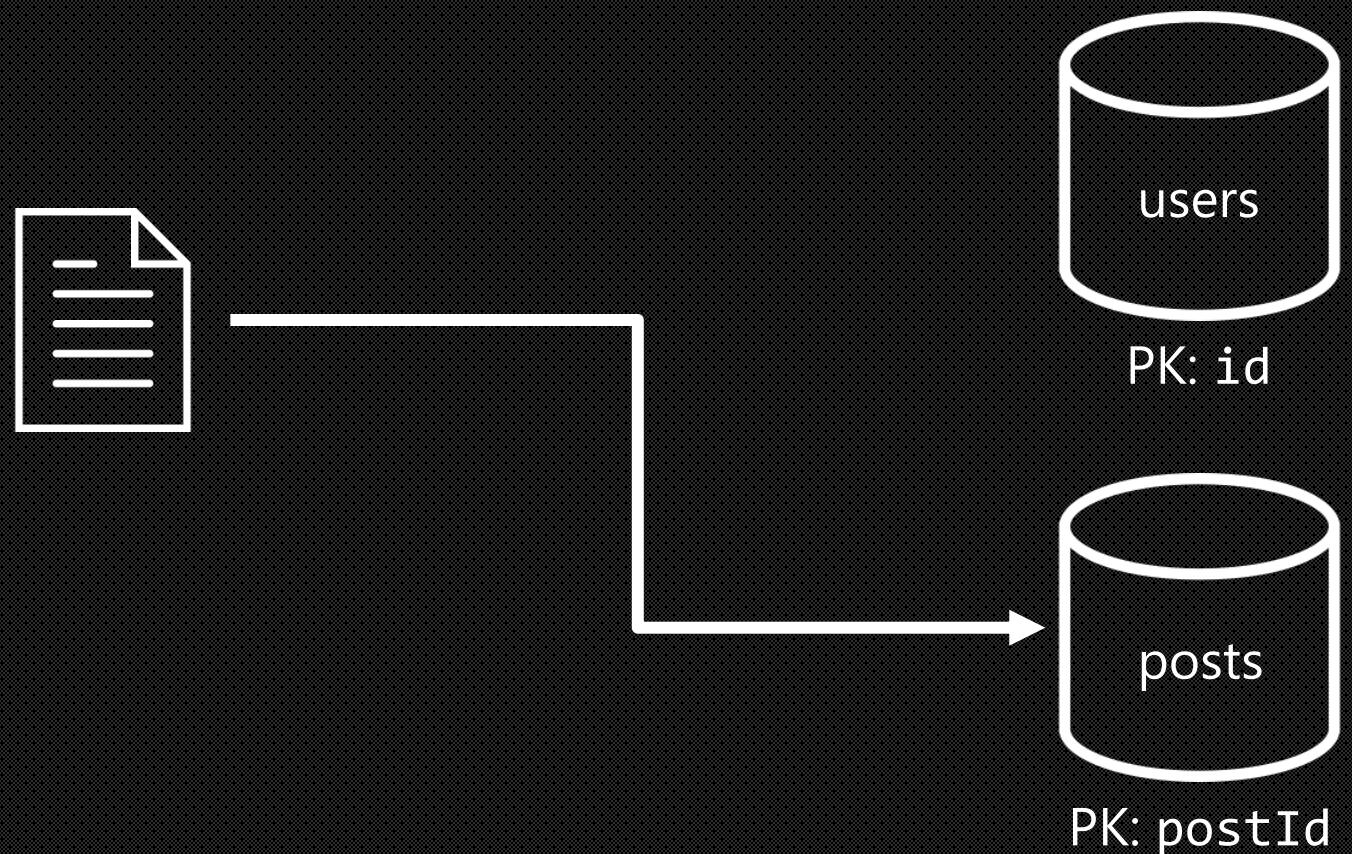
Q6



130 ms / 619.41 RU

Create a comment

C1	✓
Q1	✓
C2	✓
Q2	⚠
Q3	⚠
C3	✓
Q4	
C4	
Q5	
Q6	

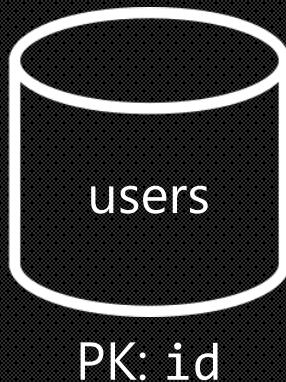


7 ms / 8.57 RU

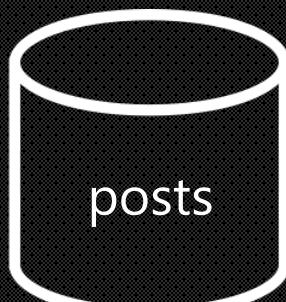
List a post's comments

C1	✓
Q1	✓
C2	✓
Q2	!
Q3	!
C3	✓
Q4	!
C4	
Q5	
Q6	

```
SELECT * FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'comment'
```



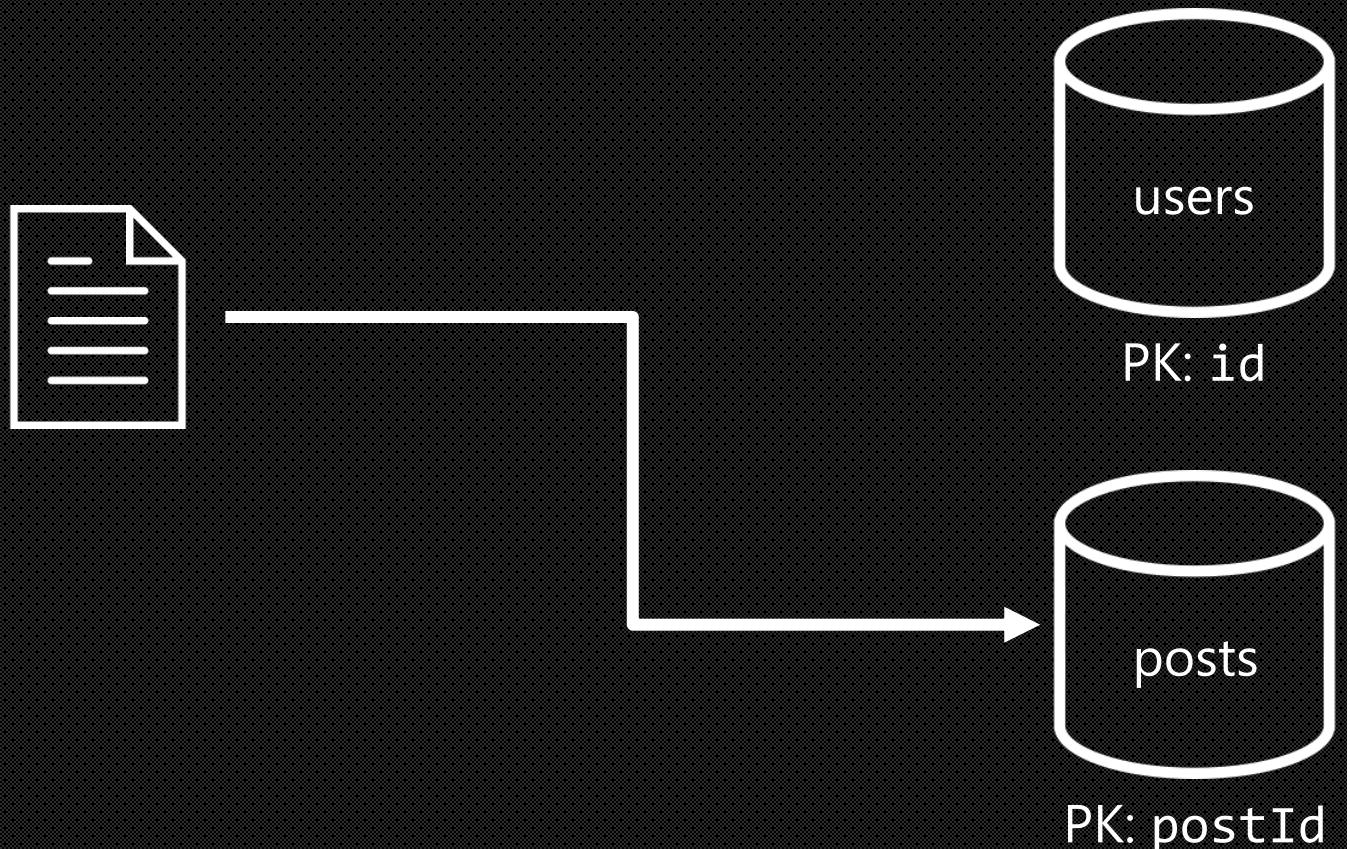
```
SELECT u.username FROM u  
WHERE u.id = '<comment-author-id>'
```



23 ms / 27.72 RU

Like a post

C1	✓
Q1	✓
C2	✓
Q2	⚠
Q3	⚠
C3	✓
Q4	⚠
C4	✓
Q5	
Q6	

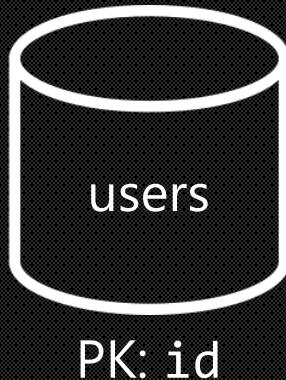


6 ms / 7.05 RU

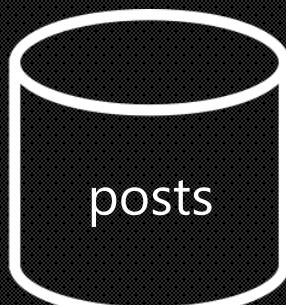
List a post's likes

C1	✓
Q1	✓
C2	✓
Q2	!
Q3	!
C3	✓
Q4	!
C4	✓
Q5	!
Q6	

```
SELECT * FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'like'
```



```
SELECT u.username FROM u  
WHERE u.id = '<liker-id>'
```



59 ms / 58.92 RU

List the x most recent posts in short form (feed)

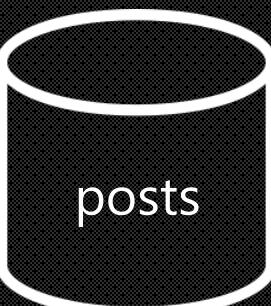
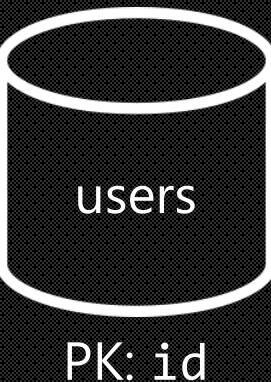
- C1 ✓
- Q1 ✓
- C2 ✓
- Q2 !
- Q3 !
- C3 ✓
- Q4 !
- C4 ✓
- Q5 !
- Q6 !

```
SELECT TOP <x> * FROM p  
WHERE p.type = 'post'  
ORDER BY p.creationDate DESC
```

```
SELECT u.username FROM u  
WHERE u.id = '<post-author-id>'
```

```
SELECT COUNT(1) FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'comment'
```

```
SELECT COUNT(1) FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'like'
```



306 ms / 2063.54 RU

Cost in RU's

C1	Create/edit a user's profile	5.71 RU
Q1	Retrieve a user's profile	2.90 RU
C2	Create/edit a post	8.76 RU
Q2	Retrieve a post	19.54 RU
Q3	List a user's post in short form	619.41 RU
C3	Create a comment	8.57 RU
Q4	List a post's comments	27.72 RU
C4	Like a post	7.05 RU
Q5	List a post's likes	58.92 RU
Q6	List the x most recent posts in short form (feed)	2,063.54 RU

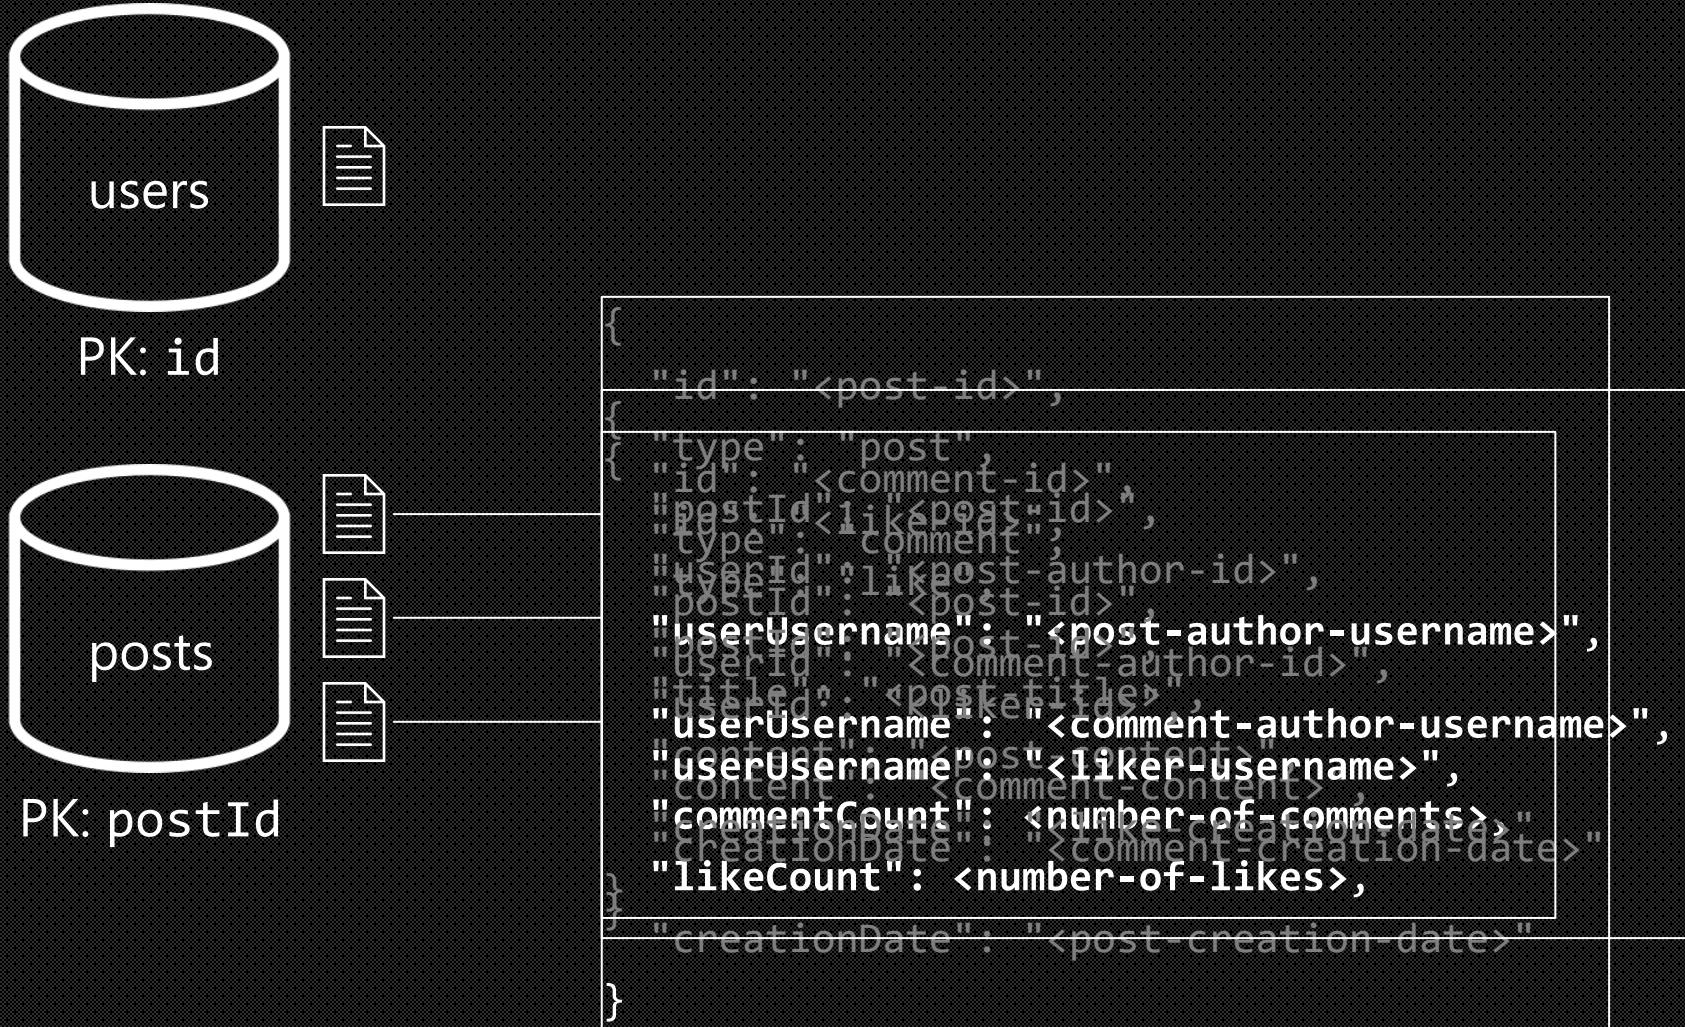
What we'll cover
today

Why Cosmos
Provisioning Cosmos
Working with Data
Data Modeling Strategies
Performance and tuning

V1: performance issues

- Having to issue multiple queries to handle a single request
- Issuing queries that don't filter on the partition key, leading to a partition scan

V2: introducing denormalization



Denormalizing within the same logical partition with stored procedures

- Written in Javascript
- Scoped to a single logical partition
- Executed as atomic transactions

Create a comment

C1

Q1

C2

Q2

Q3

C3

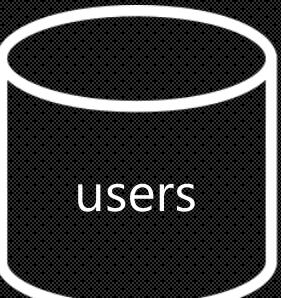
Q4

C4

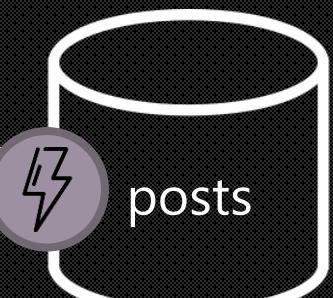
Q5

Q6

```
function createComment(postId, comment) {  
    var collection = getContext().getCollection();  
    collection.readDocument(  
        `${collection.getAltLink()}/docs/${postId}`,  
        function (err, post) {  
            post.commentCount++;  
            collection.replaceDocument(  
                post._self,  
                post,  
                function (err) {  
                    comment.postId = postId;  
                    collection.createDocument(  
                        collection.getSelfLink(),  
                        comment  
                    );  
                }  
            );  
        }  
    );  
}
```



PK: id



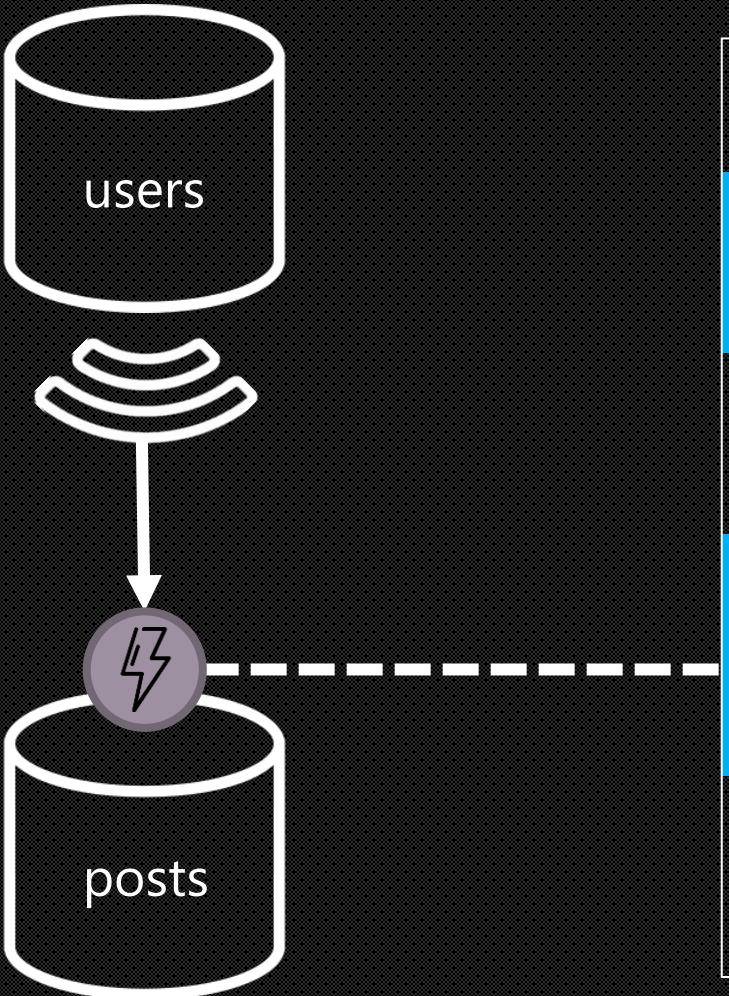
posts

PK: postId

Denormalizing across logical partitions and containers with the change feed



Denormalizing usernames into posts



```
function updateUsername(userId, username) {  
    var collection = getContext().getCollection();  
    collection.queryDocuments(  
        collection.getSelfLink(),  
        `SELECT * FROM p WHERE p.userId = '${userId}'`,  
        function (err, results) {  
            for (var i in results) {  
                var doc = results[i];  
                doc.userUsername = username;  
                collection.replaceDocument(  
                    collection.getSelfLink(),  
                    doc);  
            }  
        });  
}
```

Retrieve a post

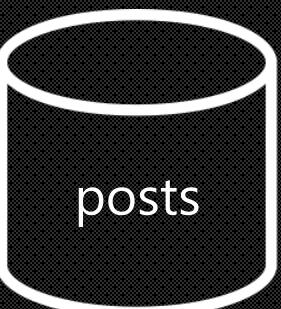
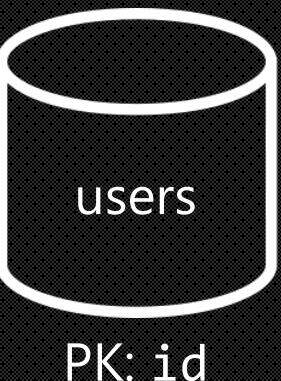
- C1 ✓
- Q1 ✓
- C2 ✓
- Q2 ✓
- Q3 !
- C3 ✓
- Q4 !
- C4 ✓
- Q5 !
- Q6 !

```
SELECT * FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'post'
```

```
SELECT u.username FROM u  
WHERE u.id = '<post-author-id>'
```

```
SELECT COUNT(1) FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'comment'
```

```
SELECT COUNT(1) FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'like'
```

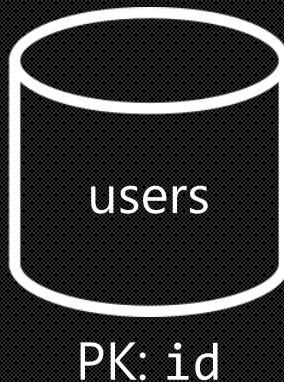


3 ms / 5.40 RU

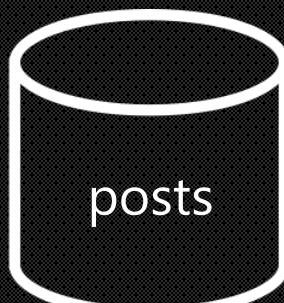
List a post's comments

- C1 ✓
- Q1 ✓
- C2 ✓
- Q2 ✓
- Q3 !
- C3 ✓
- Q4 ✓
- C4 ✓
- Q5 !
- Q6 !

```
SELECT * FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'comment'
```



```
SELECT u.username FROM u  
WHERE u.id = '<comment-author-id>'
```

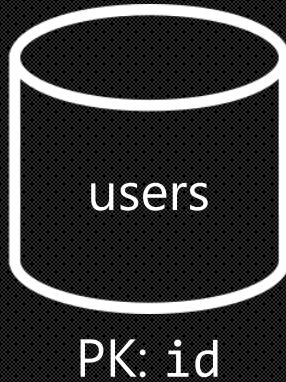


4 ms / 7.72 RU

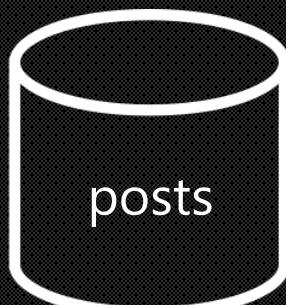
List a post's likes

- C1 ✓
- Q1 ✓
- C2 ✓
- Q2 ✓
- Q3 !
- C3 ✓
- Q4 ✓
- C4 ✓
- Q5 ✓
- Q6 !

```
SELECT * FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'like'
```



```
SELECT u.username FROM u  
WHERE u.id = '<liker-id>'
```



4 ms / 8.92 RU

List a user's posts in short form

- C1 ✓
- Q1 ✓
- C2 ✓
- Q2 ✓
- Q3 ⚠
- C3 ✓
- Q4 ✓
- C4 ✓
- Q5 ✓
- Q6 ⚠

```
SELECT * FROM p  
WHERE p.userId = '<author-id>'  
AND p.type = 'post'
```

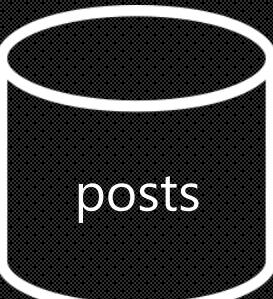
```
SELECT u.username FROM u  
WHERE u.id = '<post-author-id>'
```

```
SELECT COUNT(1) FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'comment'
```

```
SELECT COUNT(1) FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'like'
```



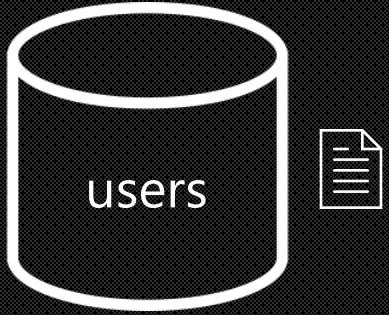
PK: id



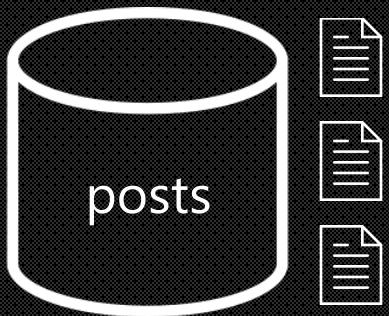
PK: postId

83 ms / 532.33 RU

V3: denormalizing entire documents

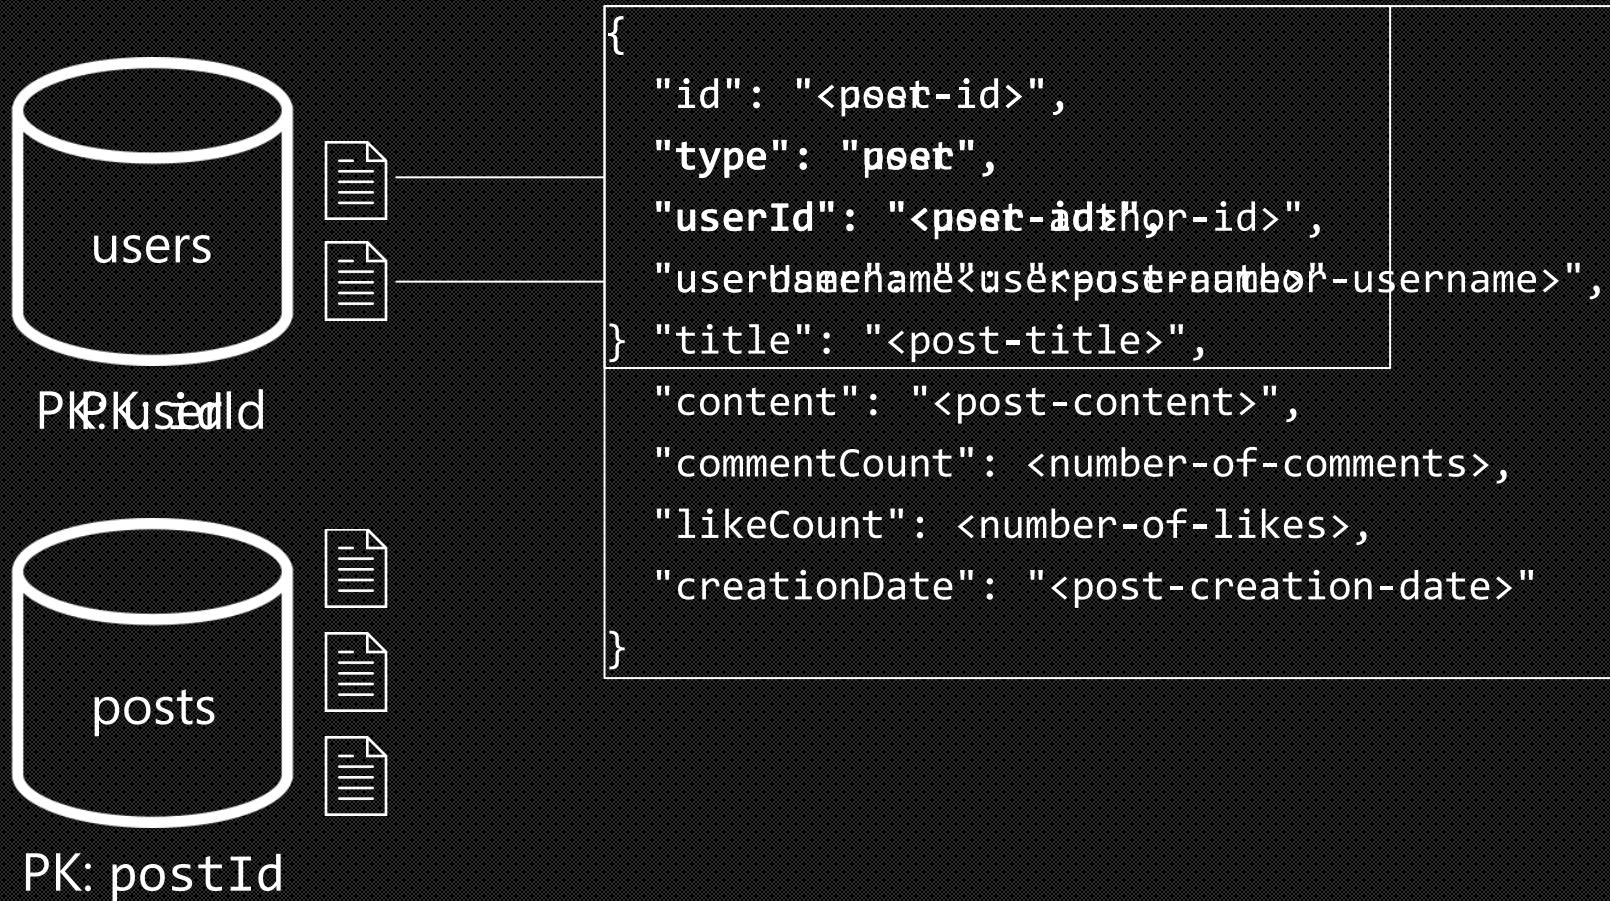


PK: id

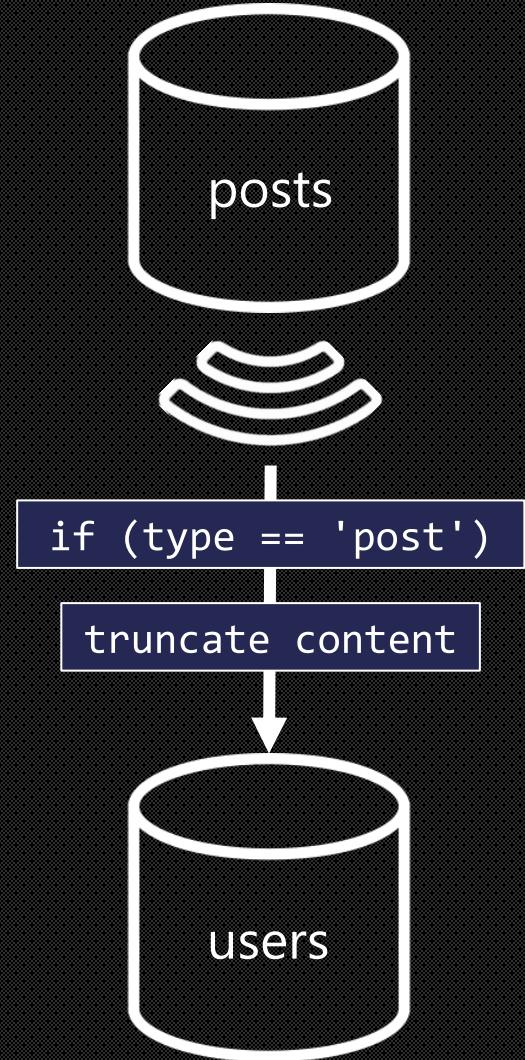


PK: postId

V3: denormalizing entire documents



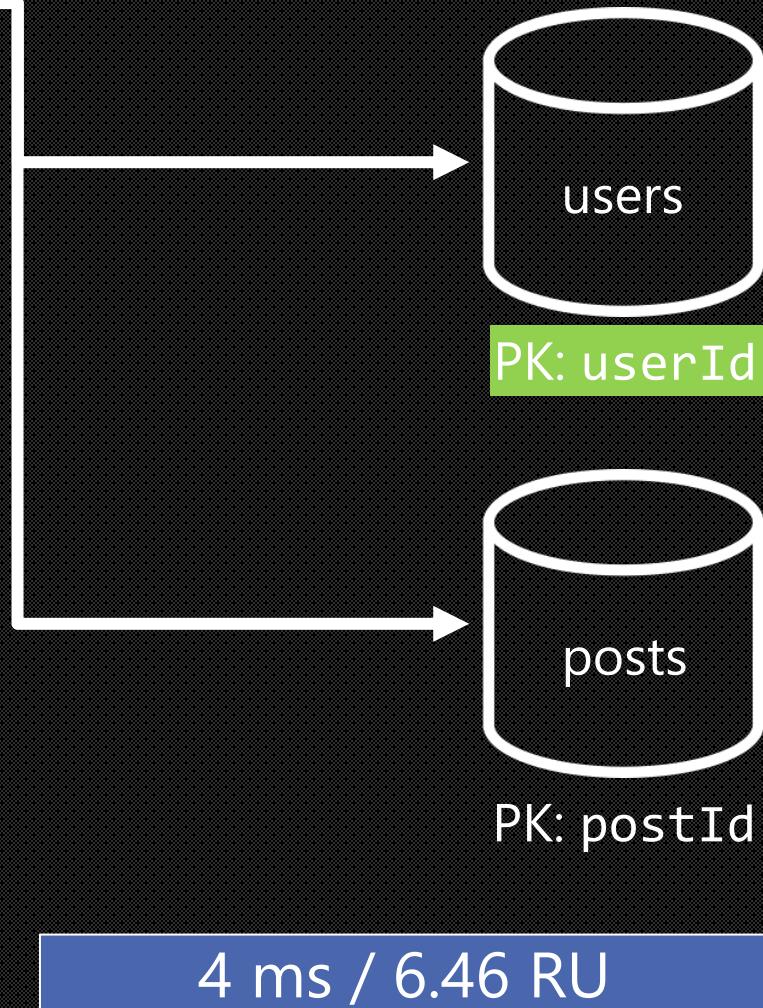
Denormalizing documents with the change feed



List a user's posts in short form

- C1 ✓
- Q1 ✓
- C2 ✓
- Q2 ✓
- Q3 ✓
- C3 ✓
- Q4 ✓
- C4 ✓
- Q5 ✓
- Q6 !

```
SELECT * FROM p  
WHERE p.userId = '<author-id>'  
AND p.type = 'post'
```



4 ms / 6.46 RU

List the x most recent posts in short form (feed)

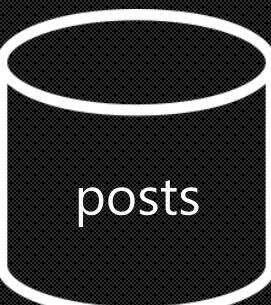
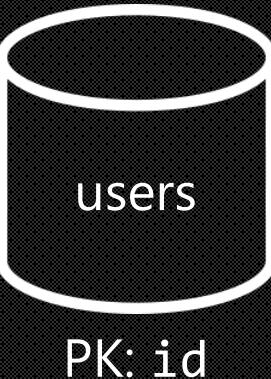
- C1 ✓
- Q1 ✓
- C2 ✓
- Q2 ✓
- Q3 ✓
- C3 ✓
- Q4 ✓
- C4 ✓
- Q5 ✓
- Q6 !

```
SELECT TOP <x> * FROM p  
WHERE p.type = 'post'  
ORDER BY p.creationDate DESC
```

```
SELECT u.username FROM u  
WHERE u.id = '<post-author-id>'
```

```
SELECT COUNT(1) FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'comment'
```

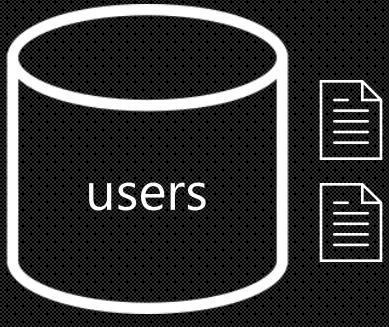
```
SELECT COUNT(1) FROM p  
WHERE p.postId = '<post-id>'  
AND p.type = 'like'
```



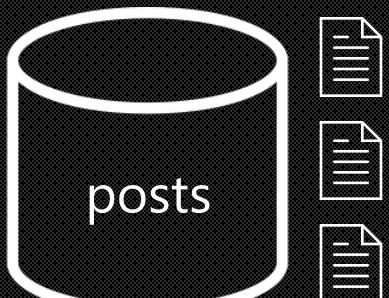
PK: postId

9 ms / 16.97 RU

V3: denormalizing entire documents

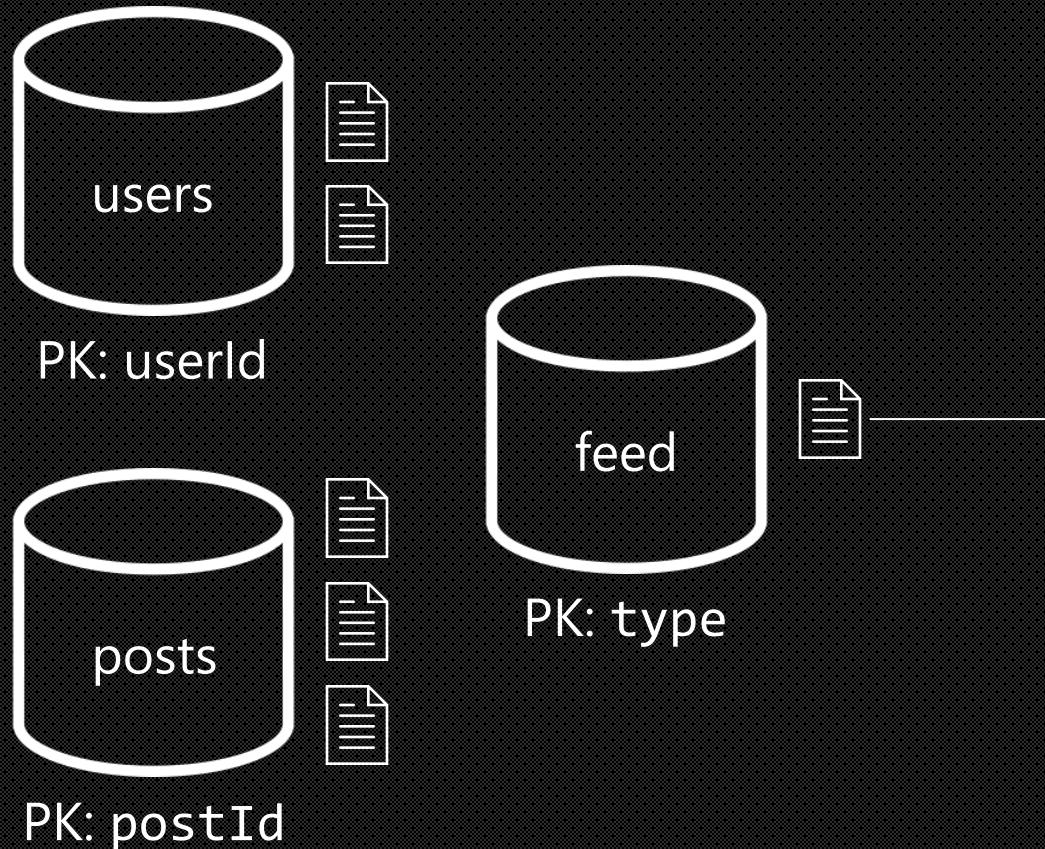


PK: userId



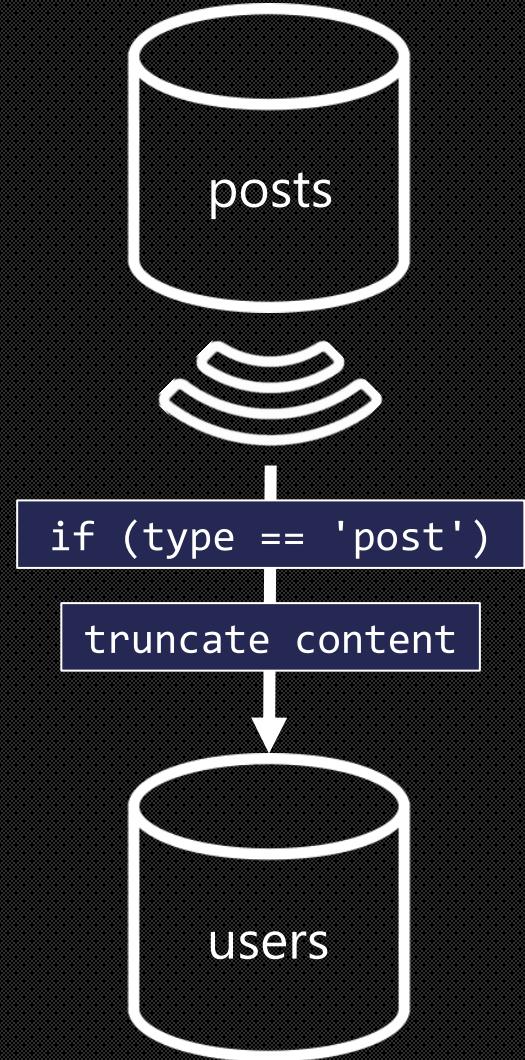
PK: postId

V3: denormalizing entire documents

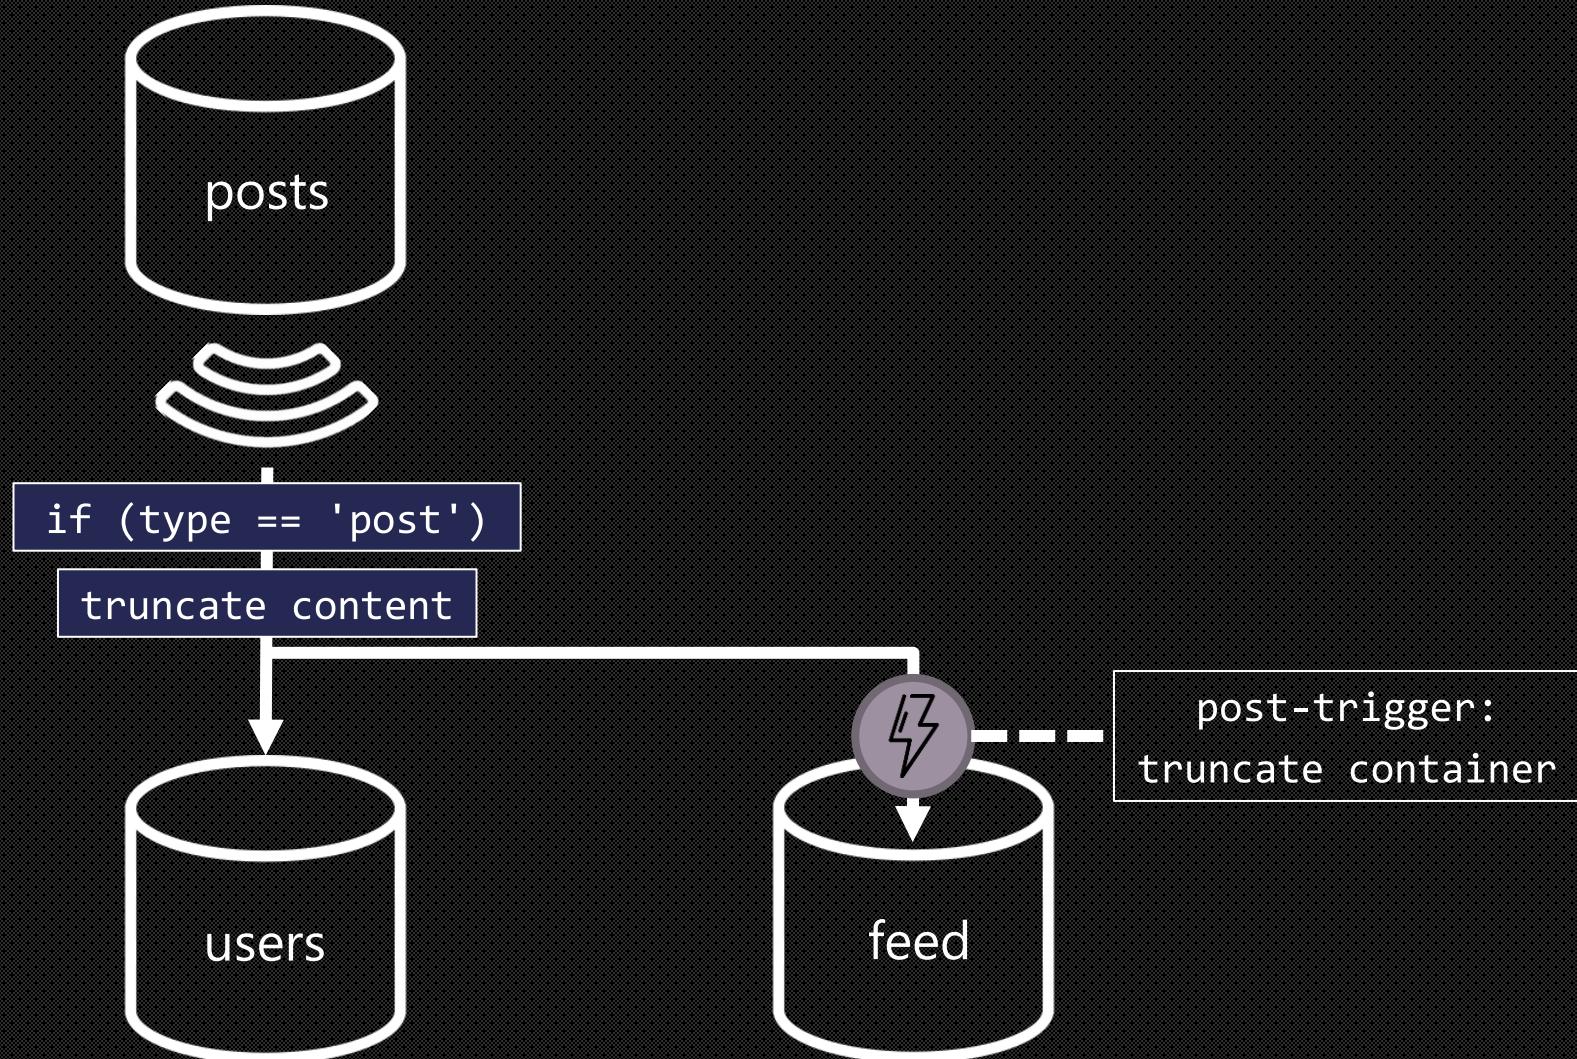


```
{  
  "id": "<post-id>",  
  "type": "post",  
  "userId": "<post-author-id>",  
  "userUsername": "<post-author-username>",  
  "title": "<post-title>",  
  "content": "<post-content>",  
  "commentCount": <number-of-comments>,  
  "likeCount": <number-of-likes>,  
  "creationDate": "<post-creation-date>"  
}
```

Denormalizing documents with the change feed



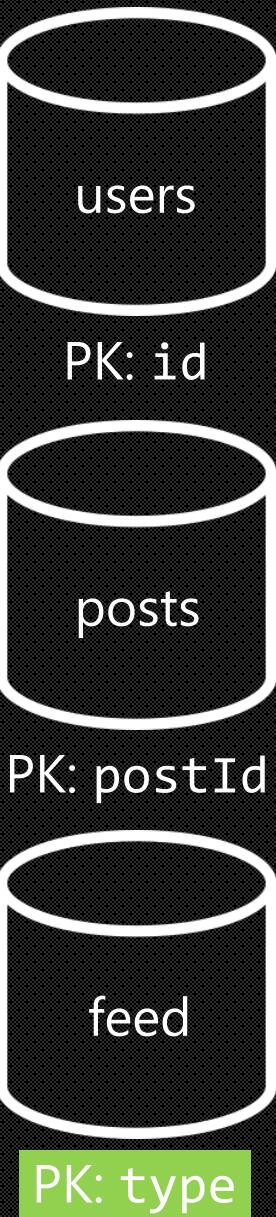
Denormalizing documents with the change feed



List the x most recent posts in short form (feed)

- C1 ✓
- Q1 ✓
- C2 ✓
- Q2 ✓
- Q3 ✓
- C3 ✓
- Q4 ✓
- C4 ✓
- Q5 ✓
- Q6 ✓

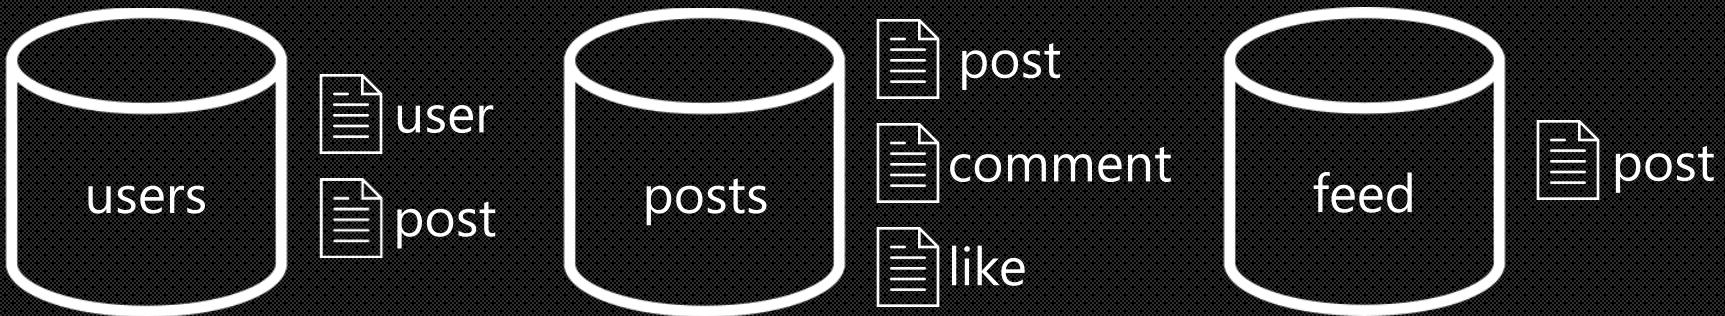
```
SELECT TOP <x> * FROM p  
WHERE p.type = 'post'  
ORDER BY p.creationDate DESC
```



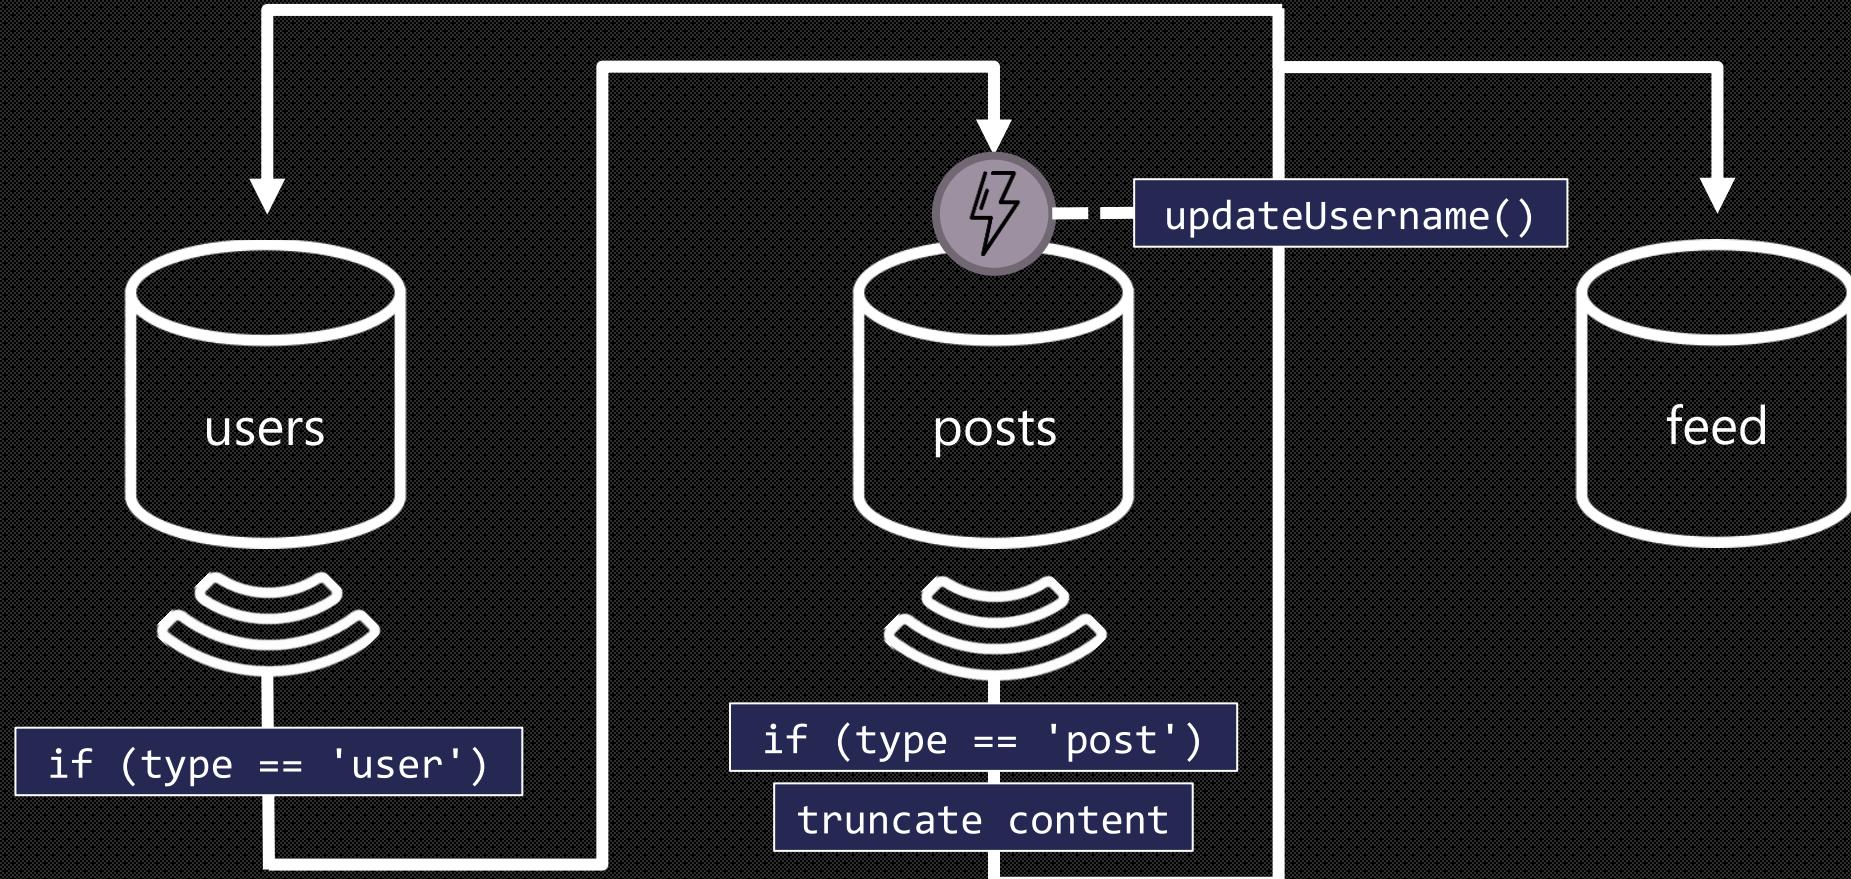
9 ms / 16.97 RU

PK: type

Our final design



Our final design



Denormalization can be added later

- V1 was not that bad!
- You can add denormalization later
 - The **change feed is persistent**
 - Run a one-time catch-up operation
- **Monitor** your database performance and setup **alerts**

	V1
C1	7 ms / 5.71 RU
Q1	3 ms / 2.90 RU
C2	9 ms / 8.76 RU
Q2	9 ms / 19.54 RU
Q3	130 ms / 619.41 RU
C3	7 ms / 8.57 RU
Q4	23 ms / 27.72 RU
C4	6 ms / 7.05 RU
Q5	59 ms / 58.92 RU
Q6	306 ms / 2063.54 RU

An overall look at performance

	V1	V2	V3
C1	7 ms / 5.71 RU	7 ms / 5.71 RU	7 ms / 5.71 RU
Q1	3 ms / 2.90 RU	3 ms / 2.90 RU	3 ms / 2.90 RU
C2	9 ms / 8.76 RU	9 ms / 8.76 RU	9 ms / 8.76 RU
Q2	9 ms / 19.54 RU	3 ms / 5.11 RU	3 ms / 5.11 RU
Q3	130 ms / 619.41 RU	28 ms / 201.54 RU	4 ms / 6.46 RU
C3	7 ms / 8.57 RU	7 ms / 15.27 RU	7 ms / 15.27 RU
Q4	23 ms / 27.72 RU	4 ms / 7.72 RU	4 ms / 7.72 RU
C4	6 ms / 7.05 RU	7 ms / 14.67 RU	7 ms / 14.67 RU
Q5	59 ms / 58.92 RU	4 ms / 8.92 RU	4 ms / 8.92 RU
Q6	306 ms / 2063.54 RU	83 ms / 532.33 RU	9 ms / 16.97 RU

What about the cost of denormalization?

- We optimized for a read-heavy scenario
 - Reads >> writes
- Q6 went from 2,000+ to 17 RUs (120x optimization)
 - Saved 1920 RUs by denormalizing posts at ~10 RUs / document
- All depends on your specific scenario and access patterns

What else do you get?

- Multi-model APIs
- Flexibility of the change feed
- Turnkey global distribution
- Industry-leading SLA

Most importantly: what's our scalability?

- V3: limitless!
- Performance will ***not*** depend on
 - number of documents
 - number of users
 - number of posts
- Whether you have 10 or 10M users, query latency will always be **the same**

	V3
C1	7 ms / 5.71 RU
Q1	3 ms / 2.90 RU
C2	9 ms / 8.76 RU
Q2	3 ms / 5.11 RU
Q3	4 ms / 6.46 RU
C3	7 ms / 14.67 RU
Q4	4 ms / 7.72 RU
C4	7 ms / 14.67 RU
Q5	4 ms / 8.92 RU
Q6	9 ms / 16.97 RU

**What we'll cover
today**

Why Cosmos
Provisioning Cosmos
Working with Data
Data Modeling Strategies
Performance and tuning

Thanks to the following

- **Thomas Weiss** presentation on Cosmos Performance in 2019
- **Alpaqa Studio**
<https://alpaqastudio.com>
- **Jeff Widmer** github code example of scenario
<https://github.com/jwidmer/AzureCosmosDbBlogExample>
- **Jane Chen** - Cosmos Jupyter Notebooks
- **Mark Brown's** Cosmos TV
<https://gotcosmos.com>
- **James Serra's** Blog
<https://www.jamesserra.com>

Mike Benkovich

- Developer
- Cloud **Architect & Consultant**
- Live in **Minneapolis**
- Founder of **Imagine Technologies, Inc.**
- Developing Courses for **LinkedIn Learning**
- Blog www.benkoTIPS.com
- Follow **@mbenko** on **Twitter**
- Send me **Feedback!** mike@benkotips.com



Mike Benkovich

Enterprise Cloud Architect,
Consultant, Developer Tools Ev...



Going further

Questions? I'm mike@benko.com - www.benkoTIPS.com

- aka.ms/PracticalCosmosDB
- cosmosdb.com
- <https://github.com/jwidmer/AzureCosmosDbBlogExample>